

Practical Lab File for Agentic AI

Sharda School of Engineering and Technology

Name- Preksha Tomar

Sys ID- 2023387520

Semester/Year- 6th sem, 3rd Year

Faculty-In-Charge/Submitted To

Mr. Ayush Singh



Sharda University

Plot No. 32-34,
Knowledge Park III,
Greater Noida, Uttar
Pradesh 201310

Experiment 1

Aim - Fine-tuning BLIP on an Image Captioning Dataset

1. Introduction

This notebook demonstrates how to **fine-tune the BLIP (Bootstrapped Language Image Pretraining)** model on a custom **image captioning dataset** using the **Hugging Face Transformers** and **Datasets** libraries.

BLIP is a vision–language model capable of understanding images and generating meaningful textual descriptions (captions).

2. Objective

The main objectives of this project are:

- To load an image–caption dataset
- To preprocess images and text using BLIP processor
- To fine-tune the BLIP model on the dataset
- To generate captions for images after training
- To visualize generated captions with images

3. Technologies & Libraries Used

- **Python**
- **PyTorch**
- **Hugging Face Transformers**
- **Hugging Face Datasets**
- **Matplotlib**
- **BLIP (Salesforce/blip-image-captioning-base)**

4. Installation of Required Libraries

```
!pip install git+https://github.com/huggingface/transformers.git@main
```

```
!pip install -q datasets
```

Purpose:

- Installs the latest version of the **Transformers** library from GitHub.
- Installs the **datasets** library for loading image-caption datasets.

5. Dataset Loading

```
from datasets import load_dataset
```

```
dataset = load_dataset("ybelkada/football-dataset", split="train")
```

Dataset Details:

- **Name:** ybelkada/football-dataset
- **Type:** Image-caption dataset
- **Split:** Training split
- Each dataset item contains:
 - image → football-related image

- text → caption describing the image

6. Dataset Inspection

```
dataset[0]["text"]
```

Purpose:

- Displays the caption associated with the first image
- Helps understand the dataset structure

7. Model & Processor Initialization

The BLIP model and processor are used to handle both image and text data.

Key Components:

- **Processor:** Converts images and text into tensors
- **Model:** Generates captions from image embeddings

8. Custom PyTorch Dataset Class

```
class ImageCaptioningDataset(Dataset):
    def __init__(self, dataset, processor):
        self.dataset = dataset
        self.processor = processor

    def __len__(self):
        return len(self.dataset)

    def __getitem__(self, idx):
        item = self.dataset[idx]
        encoding = self.processor(
            images=item["image"],
            text=item["text"],
            padding="max_length",
            return_tensors="pt"
        )
        encoding = {k: v.squeeze() for k, v in encoding.items()}
        return encoding
```

Explanation:

- Converts dataset into a format suitable for PyTorch training
- Handles:
 - Image preprocessing
 - Caption tokenization
- Removes extra batch dimensions for training compatibility

9. DataLoader Creation

The dataset is passed into a **DataLoader** for batch processing during training.

Benefits:

- Efficient batching

- Shuffling of training data
- GPU-friendly execution

10. Model Training (Fine-Tuning)

- The BLIP model is trained using:
 - Image pixel values
 - Corresponding captions
- Loss is calculated between generated and actual captions
- Optimizer updates model weights

(Exact optimizer and training loop are handled internally in the notebook)

11. Caption Generation (Inference)

```
inputs = processor(images=image, return_tensors="pt").to(device)
pixel_values = inputs.pixel_values
```

```
generated_ids = model.generate(pixel_values=pixel_values, max_length=50)
generated_caption = processor.batch_decode(
    generated_ids, skip_special_tokens=True
)[0]
```

Purpose:

- Generates captions for unseen images
- Uses trained BLIP model for prediction

12. Visualization of Results

```
import matplotlib.pyplot as plt
```

```
fig = plt.figure(figsize=(18, 14))
```

```
for i, example in enumerate(dataset):
    image = example["image"]
    ...
    plt.imshow(image)
    plt.axis("off")
    plt.title(f"Generated caption: {generated_caption}")
```

Output:

- Displays images along with generated captions
- Helps visually evaluate model performance

13. Results

- The fine-tuned BLIP model successfully generates meaningful captions
- Captions align well with football-related image content
- Visualization confirms reasonable accuracy

14. Conclusion

This notebook successfully demonstrates:

- ✓ Fine-tuning a vision-language model
- ✓ Using BLIP for image captioning
- ✓ Handling multimodal datasets
- ✓ Generating and visualizing captions

BLIP proves to be an effective model for image captioning tasks with minimal fine-tuning.

15. Future Enhancements

- Train on a larger and more diverse dataset
- Improve captions using beam search
- Deploy the model as a web application
- Add BLEU / ROUGE evaluation metrics